

CS061 - Lab 03

1 High Level Description

Today's lab will cover some new LC3 instructions, introducing the marvels of memory addressing modes, and some more I/O (input/output),

2 Goals for This Week

1. New LC3 instructions using memory addressing modes direct, indirect, and relative.
2. Exercises 1 – 4:
Practice using the three memory addressing modes direct, indirect & relative;
Building a simple loop using the conditional branch instruction; and
Simple input/output using BIOS routines.
3. So what now??

3.1 New LC3 Instructions

As you know, in the LC3, values can be stored in system memory (RAM), or in the eight general purpose registers named R0 through R7, or in several special purpose registers (*more about this later on*). It might not be obvious yet, but ***all*** microprocessor instructions involve getting values from some of those places, manipulating them, and writing the result to another place - so we need a way to describe this process.

From now on, we will use a standard notation to describe all such actions called **Register Transfer Notation (RTN)**:

- (Rn) means “the contents of Register n” - i.e the value currently stored in that register.
- Mem[xnnnn] means “the contents of memory address xnnnn”
- \leftarrow means “write the ***value*** on the right of the arrow to the ***location*** on the left” (*the location on the left may be either a memory address or a local register*), so:

$R3 \leftarrow \text{Mem}[\text{x3042}]$

means copy the contents of memory location x3042 to Register 3.

$\text{Mem}[\text{x3120}] \leftarrow (R6)$

means copy the contents of Register 6 to the memory location x3120

$R3 \leftarrow (R1) + (R2)$

means add the contents of R1 and R2, and write the result to R3.

Also, as we saw last week, a label is an alias for a memory location (an address), so:

$\text{Mem}[\text{myAnswer}] \leftarrow (R4)$

means copy the contents of R4 to the memory location corresponding to the label myAnswer.

New LC3 instructions this week: ST, LDI, STI, LDR, STR, Trap

Note: further details on all LC-3 instructions can be found [here](#) and in Appendix A of the text.

The **ST (Store)** instruction - see the [ST tutorial](#) (Piazza Resources -> LC3 Instructions)

This is simply the opposite of the LD instruction: it stores the contents of a register directly into a location in memory specified by a label (*overwriting any previous content, of course*).

The **LDI (Load Indirect)** instruction - see the [LDI tutorial](#)

This is similar to the LD instruction, except that there is one level of indirection:

Some background: As you know already, the instruction

$\text{LD } R1, \text{ ADDR_1}$

gets Mem[ADDR_1] (*i.e. the value stored at the label ADDR_1 - let's say it is x3100*) directly, and loads (writes) that value into R1:

$R1 \leftarrow \text{Mem}[\text{ADDR_1}] \Rightarrow \underline{R1 \leftarrow \text{x3100}}$

Note that the label ADDR_1 must be closer than +/- 256 locations from the instruction: if you attempt to LD or ST with a label further away than that, you will get an assembly error:

"Instruction references label that cannot be represented in a 9-bit signed PC offset" (see [here](#)).

So we need a mode that gets around that limitation - in fact, we have two such modes.

The "Load Indirect" instruction

`LDI R1, ADDR_1`

instead, loads `Mem[Mem[ADDR_1]]` into R1:

That is, it first fetches `Mem[ADDR_1]` (*let's use the same value as above, x3100*).

Then it uses that value as a **new address**, and fetches `Mem[x3100]` (*i.e. the value stored at x3100 – let's say it is #12*), and finally loads that value into R1:

$$R1 \leftarrow \text{Mem}[\text{Mem}[\text{ADDR_1}]] \quad \text{i.e.} \quad R1 \leftarrow \text{Mem}[\text{x3100}] \quad \text{i.e.} \quad \underline{R1 \leftarrow \#12}$$

This round-about process is called “**indirection**”, and the label is called a “**pointer**” (*sound familiar?*), and it allows to get around the +/- 256 line limitation

The **STI (Store Indirect)** instruction - see the [STI tutorial](#)

This works just like LDI, only in the opposite direction.

It takes the value from a register and copies (stores) it to memory at the address given by `Mem[someLabel]` (*overwriting previous content*).

$$\text{Mem}[\text{Mem}[\text{ADDR_1}]] \leftarrow (R1) \quad \Rightarrow \quad \underline{\text{x3100} \leftarrow (R1)} \quad (\text{i.e. the value stored in R1})$$

The **LDR (Load Relative)** instruction - see the [LDR tutorial](#)

This achieves the indirection like LDI, except that it uses a *register* as the pointer rather than a memory location:

The instruction

`LDR R1, R5, #0` ; For this course, we will always set the offset to 0

loads the value from `Mem[(R5) + offset 0]` into R1: (*let's say (R5) == x3100, and Mem[x3100] is #12*)

$$R1 \leftarrow \text{Mem}[(R5)] \quad \text{i.e.} \quad R1 \leftarrow \text{Mem}[\text{x3100}] \quad \text{i.e.} \quad \underline{R1 \leftarrow \#12}$$

The **STR (Store Relative)** instruction - see the [STR tutorial](#)

This works just like LDR, only in the opposite direction.

In the example, the instruction stores the value from R1 into the address stored in R5 (x3007):

The instruction

`STR R1, R5, #0`

copies the value from R1 (*let's say #12*) into the address stored in (R5): (*let's say (R5) == x3100, as above*)

$$\text{Mem}[(R5)] \leftarrow (R1) \quad \text{i.e.} \quad \underline{\text{x3100} \leftarrow \#12}$$

The **Trap** instruction - see Table 1 below and the [TRAP tutorial](#) and p. 543 in the text. **Trap** instructions are actually calls to a set of **BIOS subroutines** (what C++ calls “functions”). You saw one of them (PUTS) last week. You may use either the “TRAP xnn” form or the alias in your code.

Table 1: Trap instructions

Invocation	Alias	Effect
TRAP x20	GETC	R0 ← one character of input (ascii) from the keyboard (no echo)
TRAP x21	OUT	print (R0) to the screen as an ascii character
TRAP x22	PUTS	Jump to memory location (R0) and print the contents of that and each succeeding memory location until a 0 is encountered (used for printing strings)
TRAP x23	IN	<i>You will never use this instruction!</i> Prompt the user to input a character; get the input from the keyboard, and echo it to the console
TRAP x25	HALT	Terminates the program

3.2 Exercises 1 – 4

REMINDER: Always open the Text Window of your simpl LC-3 emulator!

(by checking the "Text Window" box at the bottom as soon as you open simpl)

Exercise 01: Direct memory addressing mode

Write an assembly language program that uses .FILL pseudo-ops to load the values #65 (decimal 65) and x41 (hexadecimal 41) into two memory locations with the labels DEC_65 and HEX_41 respectively.

Look up an ASCII table (use the back of your text or [this table](#) to see what these values represent when interpreted as characters rather than numbers.

Then use the **LD** instruction to load these two values into registers R3 and R4 respectively. Run the program, and inspect the registers to make sure it did what you expected.

Exercise 02: Indirect memory addressing mode

Sometimes, the data we need is located in some remote region of memory.

Take the program from exercise 01 (*use your lab03_ex2.asm file*) and replace the data stored at DEC_65 and HEX_41 with two far away **addresses** such as x4000 and x4001 respectively, and change the labels to DEC_65_PTR and HEX_41_PTR (*because your labels are now pointers*).

You can load the data into the new locations by putting this at the end of your current data block, before the .END pseudo-op:

```
;; Remote data
                .orig x4000
NEW_DEC_65     .FILL #65
NEW_HEX_41     .FILL x41
```

But now we have a problem!

As we have already seen, the LC3 Direct memory addressing mode (LD, ST instructions) only works with labels that are within +/- #256 memory locations of the instruction (see [here](#) for full details) – so our new data locations are too far away from the code to be accessed with LD and ST, even though we have provided new labels for them.

Try using the LD instruction with these new labels and see what happens!

But never fear – the **indirect** and **relative** addressing modes will come to the rescue!

Replace the LD instruction with one using **LDI** to load the data into R3 and R4

*(Hint: use the **_PTR** labels, not the **NEW_** ones – in fact you can now remove those, they are of no use!)*

Next, add code to increment the values in R3 and R4 by 1.

(What ascii characters do they correspond to now?)

Finally, store the incremented values back into addresses x4000 and x4001 using **STI**.

Again, inspect the registers to make sure it worked as expected!

Exercise 03: Relative memory addressing mode

Use your lab03_ex3.asm file

Start with the same setup as in exercise 02: You have two labeled locations (“pointers”), which contain two “remote” memory addresses; your actual data is stored at those remote addresses.

Directly load (**LD**) the values of the pointers into R5 and R6 respectively.

Now, use the relative memory addressing mode (**LDR**) to load the remote data into R3 and R4, using R5 and R6 as “Base Registers” - i.e. the registers that hold the memory pointers.

Perform the same manipulation as in exercise 02 – i.e. increment the values in R3 & R4, then store those incremented values back into x4000 and x4001, this time using **STR** (*inspect your registers to confirm, as always*).

SUCCESS!!

You have now used the two memory addressing modes indirect (**LDI**, **STI**) and relative (**LDR**, **STR**) to accomplish exactly the same goal – each uses indirection (“pointers”) to access memory locations that were too far away to be reached by LD:

The Indirect memory addressing mode uses a label (aliased memory location) as a pointer.

The Relative memory addressing mode uses a register containing a memory location as a pointer.

(The downside of both is that they require one more memory read than direct addressing).

Exercise 04: Loops

Review the workings of the NZP condition codes and the BR instruction [here](#).

Use the conditional branch instruction (**BR**) to construct a counter-controlled loop.

Hard-code (*i.e.* use `.FILL`) local data values `x61` and `x1A`, and load them into `R0` and `R1` respectively. Inside a loop, output to console the contents of `R0` (Trap `x21`, or `OUT`), then immediately increment `R0` by 1.

Use `R1` as the loop counter – i.e. the loop should be executed exactly `x1A` times.

Think carefully about how and when to terminate the loop: do you use `BRn` or `BRnz` or `BRz` or `BRzp` or ??

What does this loop do? Can you figure it out **before** you run it?

3.3 Submission

Add, commit, and push your `lab03_ex1.asm` through `lab03_ex4.asm` files to your lab 3 GitHub repo, and demo to your TA.

4 So what do I know now?

You should now ...

- Be totally familiar with the way of describing instructions, Register Transfer Notation (RTN)
This will become fundamental to our understanding of how a microprocessor functions, which is the ultimate purpose of the entire course – so make sure you get the hang of it!
- understand the three LC3 memory addressing modes – direct, indirect, and relative
- master the `LD/ST`, `LDI/STI`, `LDR/STR` instructions, and the concept of indirection (pointers) – specifically, how to load from and store to locations in memory that are too far away from your code to be reached by the direct addressing mode (`LD/ST`).
- know how to build a simple counter-controlled loop
- Be able to use all the bios routines – i.e. the Trap instructions that manage console i/o: `GETC`, `OUT`, and `PUTS`